



Introduction to Machine Learning


WSS ML Workshop

Hosein Hasani




WSS 2024

Clustering: K-means Algorithm




Dimensionality Reduction

- A technique to find a lower-dimensional representation of data features that preserves some of its properties.




Dimensionality Reduction

- A technique to find a lower-dimensional representation of data features that preserves some of its properties.
- Motivations:
 - Computation
 - Visualization
 - Feature extraction




Dimensionality Reduction

- A technique to find a lower-dimensional representation of data features that preserves some of its properties.
- Case Study:
Principal Component Analysis (PCA)



Dimensionality Reduction

More sophisticated methods:



Variational Autoencoder (VAE)

Unsupervised Learning (Recap)

Unsupervised learning

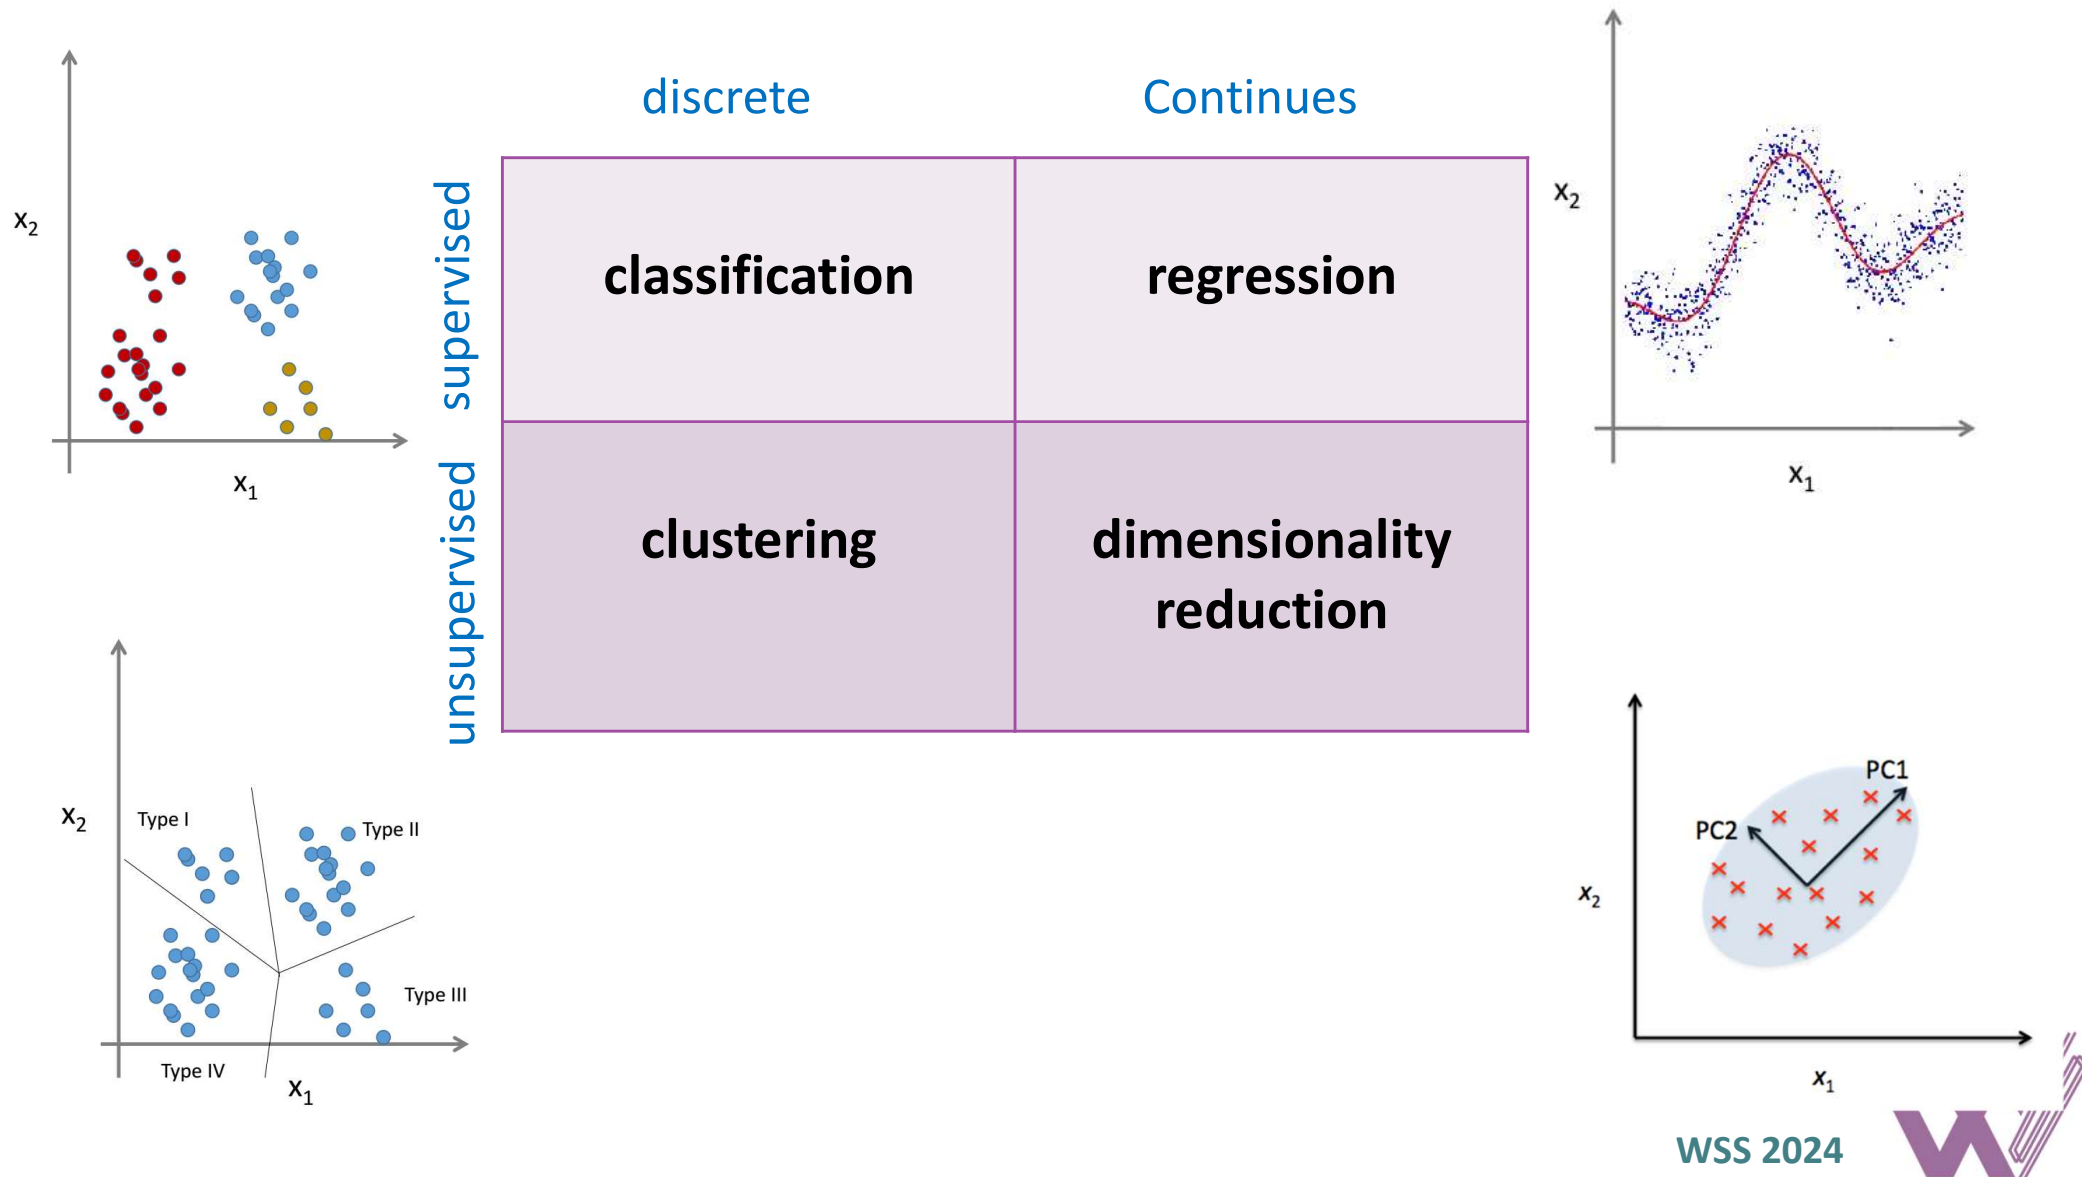
Given: Training set

$$D = \{(x^{(i)})\}_{i=1}^N$$

Goal: Revealing structure in the observed data and finding groups or intrinsic structures in the data

- Main Approaches:
 - Density Estimation
 - Generative Modelling
 - Clustering
 - Dimensionality Reduction

Supervised Learning vs. Unsupervised Learning




Outline

- Introduction and Motivation
- ML Problems
 - Supervised Learning
 - Unsupervised Learning
 - Reinforcement Learning
- Loss Function and Optimization
- Generalization and Overfitting

Reinforcement Learning


- Most natural way of learning
- Examples:
 - Baby movement
 - Investment

Agent (Baby)




Reinforcement Learning

- Sequential decision making with (possibly delayed) rewards
- An agent learns appropriate actions (policy) based on environment feedbacks to maximize reward.



Reinforcement Learning

- Sequential decision making with (possibly delayed) rewards
- Data in supervised learning:
(input, label)
- Data in reinforcement learning:
(input, some output, a grade of reward for this output)



Reinforcement Learning

- State: Agent's observation from the world
- Environment model:
 - Transition probability $p(s_{t+1}|s_t, a_t)$
 - Reward function $R(s_t, a_t, s_{t+1})$
- Policy: Mapping from states to actions
$$\pi_{\theta}: S \rightarrow A$$
- Goal: Learning an optimal policy in order to maximize its long-term reward


Multi-Armed Bandit



Multiple bandits with unknown average rewards

Multi-Armed Bandit

- Finding the **best arm** (in the sense of expected reward) with minimum trial and error.
- Minimizing cumulative **regret**.
- **Exploration-exploitation** trade-off!




Multi-Armed Bandit

Applications:

- Online advertisement
- Recommender systems
- Clinical trials
- Mining
- Network (packet routing)

Primary ML Problems (Review)



Hypothesis Class (Recap)

- The aim of supervised learning is to find f^* (best solution) from **hypothesis space** (e. g. the set of all possible functions)
- Example:

In linear regression the hypothesis space include all possible functions with the form of

$$f(x; w_0, w_1) = w_0 + w_1 x$$



set of all possible
functions H

• f^*

Loss Function and Optimization

- **Loss Function:**

Quantifies how much undesirable is each parameter vector across the training data.

- **Optimization:**

Apply an **optimization** algorithm that finds the parameters that minimize the loss function.

Steps of Learning Procedure


Typical steps of solving (supervised) learning problems:

- Select the **hypothesis space**:
- Define a **loss function** that quantifies how much undesirable is each parameter vector across the training data.
- Apply an **optimization** algorithm that efficiently finds the parameters that minimize the loss function.
- **Evaluate** the obtained model.

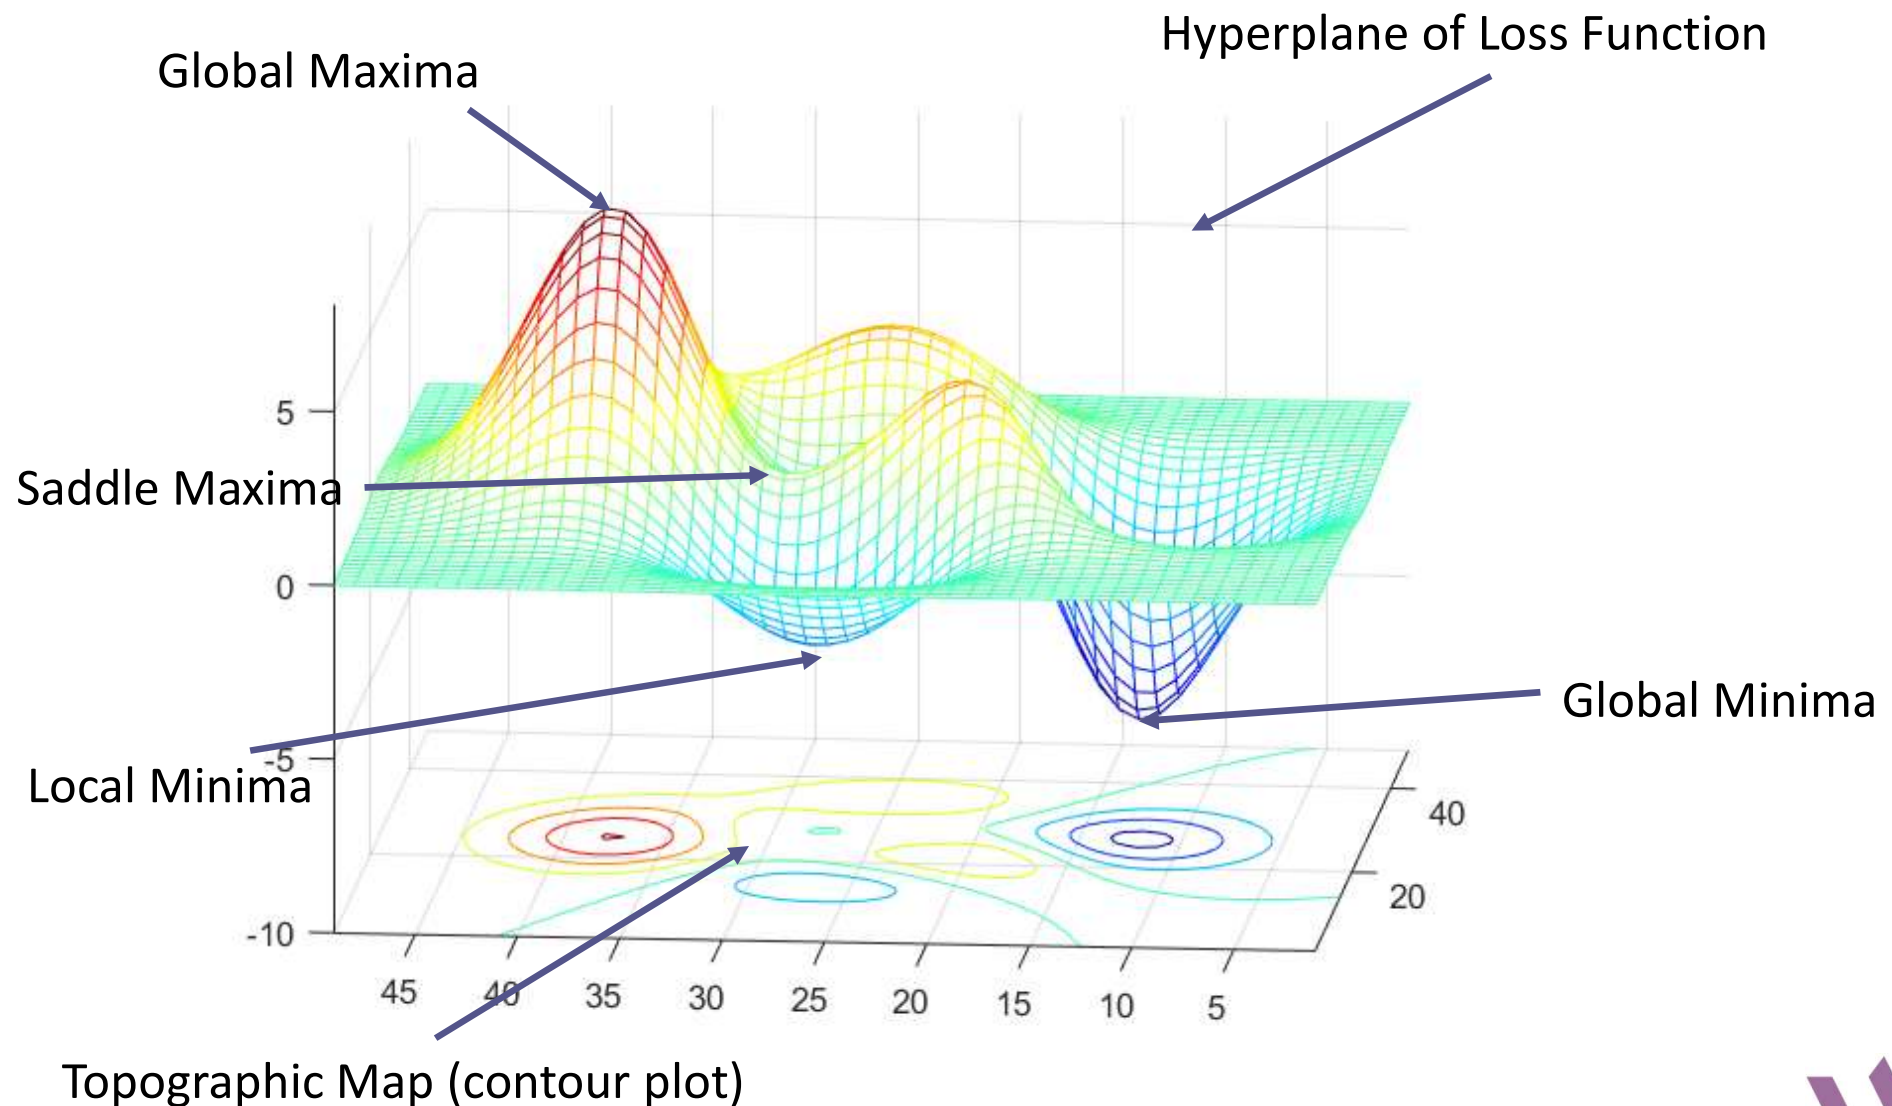
Loss Function

- Error: The **difference** between the actual outputs (e.g. **ground truth**) and the predicted outputs.
- The function that is used to compute this error is known as Loss Function $L(.)$ or $J(.)$.
- Generally, consists of **empirical risk term** and **regularization term**.

Loss Function: Loss Landscape




Loss Function: Loss Landscape




Loss Functions: Negation Problem


- Consider a threshold logic unit with a single input and training examples for the negation:




- Error of computing the negation w.r.t. the θ and w :



error for $x = 0$



error for $x = 1$



sum of errors

$$L = \sum_i L_i$$

Loss Functions in Supervised Learning

- Regression Losses
 - Mean Square Error (L2 Loss)
- Classification Losses
 - Hinge Loss (Multi class SVM)
 - Cross Entropy Loss

Mean Squared Error


Mean Squared Error (MSE) loss function is widely used in **regression** problems.

$$J(w) = \sum_{i=1}^N (y^{(i)} - \underbrace{w^T x^{(i)}}_{\widehat{y^{(i)}}})^2$$

Mean Squared Error

Mean Squared Error (MSE) loss function is widely used in **regression** problems.

$$J(w) = \sum_{i=1}^N (y^{(i)} - \underbrace{w^T x^{(i)}}_{\hat{y}^{(i)}})^2$$



Mean Squared Error

Mean Squared Error (MSE) loss function is widely used in **regression** problems.

$$J(w) = \sum_{i=1}^N (y^{(i)} - \underbrace{w^T x^{(i)}}_{\widehat{y^{(i)}}})^2$$

Goal: Find w^* which minimizes $J(w)$:

$$w^* = \operatorname{argmin}_w J(w)$$

Mean Squared Error

Mean Squared Error (MSE) loss function is widely used in **regression** problems.

$$J(w) = \sum_{i=1}^N (y^{(i)} - \underbrace{w^T x^{(i)}}_{\widehat{y}^{(i)}})^2$$

Goal: Find w^* which minimizes $J(w)$:

$$w^* = \operatorname{argmin}_w J(w)$$


Optimization

Loss Functions in Supervised Learning

- Regression Losses
 - Mean Square Error (L2 Loss)
- Classification Losses
 - Hinge Loss (Multi class SVM)
 - Cross Entropy Loss

Hinge Loss: SVM Classifier


Support Vector Machine (SVM)



[Mohri]

Hinge Loss: SVM Classifier

Support Vector Machine (SVM)




Misclassification Error: $- \text{sign}(y^{(i)}[w \cdot x^{(i)} + b])$

[Mohri]

Hinge Loss: SVM Classifier

Support Vector Machine (SVM)




Misclassification Error: $- \text{sign}(y^{(i)}[w \cdot x^{(i)} + b])$

[Mohri]

Hinge Loss: SVM Classifier

Support Vector Machine (SVM)



Misclassification Error: $- \text{sign}(y^{(i)}[w \cdot x^{(i)} + b])$

[Mohri]

Cross Entropy Loss

- Cross-entropy:

$$H(q, p) = - \sum_x q(x) \log p(x)$$

Cross Entropy Loss

- Cross-entropy:

$$H(q, p) = - \sum_x q(x) \log p(x)$$

- Multi-class cross-entropy Loss:

$$L(\hat{y}, y) = - \sum_j y_j \log \hat{y}_j$$

Cross Entropy Loss

- Multi-class cross-entropy Loss:

$$L(\hat{y}, y) = - \sum_j y_j \log \hat{y}_j$$

predicted probability of each class!

Cross Entropy Loss: Softmax Classifier

- Multi-class cross-entropy Loss:

$$L(\hat{y}, y) = - \sum_j y_j \log \hat{y}_j$$

predicted probability of each class!

- Softmax classifier:

$$s = f(x)$$

$$\hat{y}_j = \frac{e^{s_j}}{\sum_{k=1}^C e^{s_k}}$$

Cross Entropy Loss: Softmax Classifier

- Multi-class cross-entropy Loss:

$$L(\hat{y}, y) = - \sum_j y_j \log \hat{y}_j$$

predicted probability of each class!

The true distribution (one-hot vector: $q = [0, 0, 0, \dots, 1, \dots, 0]$)

- Softmax classifier:

$$s = f(x)$$

$$\hat{y}_j = \frac{e^{s_j}}{\sum_{k=1}^C e^{s_k}}$$

Loss Functions in Supervised Learning

- Regression Losses
 - Mean Square Error (L2 Loss)
- Classification Losses
 - Hinge Loss (Multi class SVM)
 - Cross Entropy Loss